


Sigmoid Exam



Data Analysis

Raw Data:

company	rank	rank_ch	revenue	profit	num. of empl	sector	city	state	newco	ceo_fou	ceo_w	profitabl	prev_	CEO	Website	Ticker	Market Cap
Walmart	1	0	572754	13673	2300000	Retailing	Bentonville	AR	no	no	no	yes	1	C. Douglas McMillon	https://www.stock.walmart.com	WMT	352037
Amazon	2	0	469822	33364	1608000	Retailing	Seattle	WA	no	no	no	yes	2	Andrew R. Jassy	www.amazon.com	AMZN	1202717
Apple	3	0	365817	94680	154000	Technology	Cupertino	CA	no	no	no	yes	3	Timothy D. Cook	www.apple.com	AAPL	2443962
CVS Health	4	0	292111	7910	258000	Health Care	Woonsocket	RI	no	no	yes	yes	4	Karen Lynch	https://www.cvshealth.com	CVS	125204
UnitedHealth Group	5	0	287597	17285	350000	Health Care	Minnetonka	MN	no	no	no	yes	5	Andrew P. Witty	www.unitedhealthgroup.com	UNH	500468
Exxon Mobil	6	4	285640	23040	63000	Energy	Irving	TX	no	no	no	yes	10	Darren W. Woods	www.exxonmobil.com	XOM	371841
Berkshire Hathaway	7	-1	276094	89795	372000	Financials	Omaha	NE	no	no	no	yes	6	Warren E. Buffett	www.berkshirehathaway.com	BRKA	625468
Alphabet	8	1	257637	76033	156500	Technology	Mountain View	CA	no	no	no	yes	9	Sundar Pichai	https://www.abc.xyz	GOOGL	1309359
McKesson	9	-2	238228	-4539	67500	Health Care	Irving	TX	no	no	no	no	7	Brian S. Tyler	www.mckesson.com	MCK	47377
AmerisourceBergen	10	-2	213988.8	1539.9	40000	Health Care	Conshohocken	PA	no	no	no	yes	8	Steven H. Collis	www.amerisourcebergen.com	ABC	29972
Costco Wholesale	11	1	195929	5007	288000	Retailing	Issaquah	WA	no	no	no	yes	12	W. Craig Jelinek	www.costco.com	COST	230443
Cigna	12	1	174078	5365	72963	Health Care	Bloomfield	CT	no	no	no	yes	13	David Cordani	https://www.cigna.com	CI	88459
AT&T	13	-2	168864	20081	202600	Telecommunications	Dallas	TX	no	no	no	yes	11	John T. Stankey	www.att.com	T	148907
Microsoft	14	1	168088	61271	181000	Technology	Redmond	WA	no	no	no	yes	15	Satya Nadella	www.microsoft.com	MSFT	1941033
Cardinal Health	15	-1	162467	611	46827	Health Care	Dublin	OH	no	no	no	yes	14	Jason Hollar	www.cardinalhealth.com	CAH	15169
Chevron	16	11	162465	15625	42595	Energy	San Ramon	CA	no	no	no	yes	27	Michael K. Wirth	www.chevron.com	CVX	284132
Home Depot	17	1	151157	16433	490600	Retailing	Atlanta	GA	no	no	no	yes	18	Edward P. Decker	www.homedepot.com	HD	308152
Walgreens Boots Alliance	18	-2	148579	2542	258500	Food & Drug Stores	Deerfield	IL	no	no	yes	yes	16	Roz Brewer	www.walgreensbootsalliance.com	WBA	33360
Marathon Petroleum	19	13	141032	9738	17700	Energy	Findlay	OH	no	no	no	yes	32	Michael J. Hennigan	www.marathonpetroleum.com	MPC	47526
Elevance Health	20	3	138639	6104	98200	Health Care	Indianapolis	IN	no	no	yes	yes	23	Gail K. Boudreaux	www.elevancehealth.com	ELV	119923
Kroger	21	-4	137888	1655	420000	Food & Drug Stores	Cincinnati	OH	no	no	no	yes	17	W. Rodney McMullen	www.thekrogerco.com	KR	33846
Ford Motor	22	-1	136341	17937	183000	Motor Vehicles & Parts	Dearborn	MI	no	no	no	yes	21	James D. Farley Jr.	www.ford.com	F	50609
Verizon Communications	23	-3	133613	22065	118400	Telecommunications	New York	NY	no	no	no	yes	20	Hans E. Vestberg	www.verizon.com	VZ	211872
JPMorgan Chase	24	-5	127202	48334	271025	Financials	New York	NY	no	no	no	yes	19	James Dimon	www.jpmorganchase.com	JPM	336469
General Motors	25	-3	127004	10019	157000	Motor Vehicles & Parts	Detroit	MI	no	no	yes	yes	22	Mary T. Barra	www.gm.com	GM	50156
Centene	26	-2	125982	1347	72500	Health Care	St. Louis	MO	no	no	yes	yes	24	Sarah M. London	www.centene.com	CNC	53429
Meta Platforms	27	7	117929	39370	71970	Technology	Menlo Park	CA	no	yes	no	yes	34	Mark Zuckerberg	https://investor.fb.com	META	475718
Comcast	28	-2	116385	14159	189000	Telecommunications	Philadelphia	PA	no	no	no	yes	26	Brian L. Roberts	www.comcastcorporation.com	CMCSA	185069
Phillips 66	29	19	114852	1317	14000	Energy	Houston	TX	no	no	no	yes	48	Mark E. Lashier	www.phillips66.com	PSX	41091
Valero Energy	30	23	108332	930	9804	Energy	San Antonio	TX	no	no	no	yes	53	Joseph W. Gorder	www.valero.com	VLO	44376

General info


```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   company             1000 non-null   object 
1   rank                1000 non-null   int64  
2   rank_change         1000 non-null   float64
3   revenue             1000 non-null   float64
4   profit              997 non-null    float64
5   num. of employees   999 non-null    float64
6   sector              1000 non-null   object 
7   city                1000 non-null   object 
8   state               1000 non-null   object 
9   newcomer            1000 non-null   object 
10  ceo_founder         1000 non-null   object 
11  ceo_woman           1000 non-null   object 
12  profitable           1000 non-null   object 
13  prev_rank           1000 non-null   object 
14  CEO                 1000 non-null   object 
15  Website             1000 non-null   object 
16  Ticker              951 non-null    object 
17  Market Cap          969 non-null    object 
dtypes: float64(4), int64(1), object(13)
memory usage: 140.8+ KB
```


```
numerical_features = df[['revenue', 'profit', 'num. of employees', 'Market Cap']]
numerical_features_description_updated = numerical_features.describe()
```

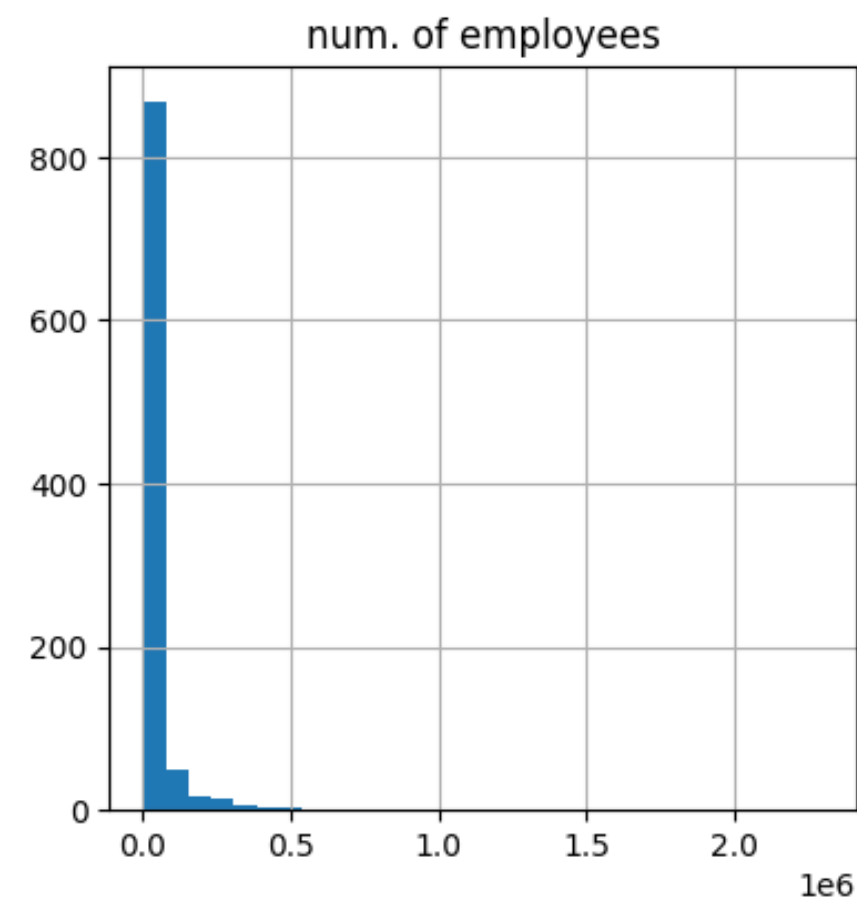


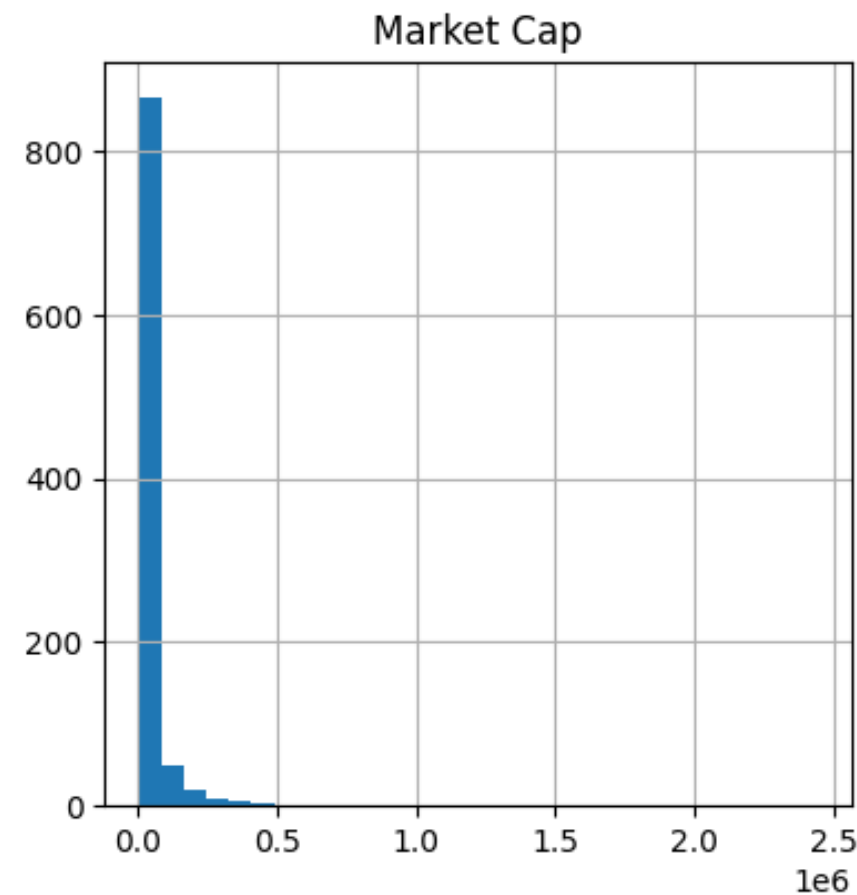
	revenue	profit	num. of employees	Market Cap
count	961.000000	958.000000	9.610000e+02	9.610000e+02
mean	17999.396462	2070.201983	3.651696e+04	4.002294e+04
std	41463.958721	6543.255875	1.063753e+05	1.309925e+05
min	2107.200000	-6520.000000	1.600000e+02	1.490000e+02
25%	3494.800000	193.500000	6.640000e+03	4.415600e+03
50%	6310.200000	573.150000	1.400000e+04	1.164160e+04
75%	14298.000000	1502.250000	3.000000e+04	3.345560e+04
max	572754.000000	94680.000000	2.300000e+06	2.443962e+06

Detecting outliers with boxplot



Distribution of Market Cap in S&P500






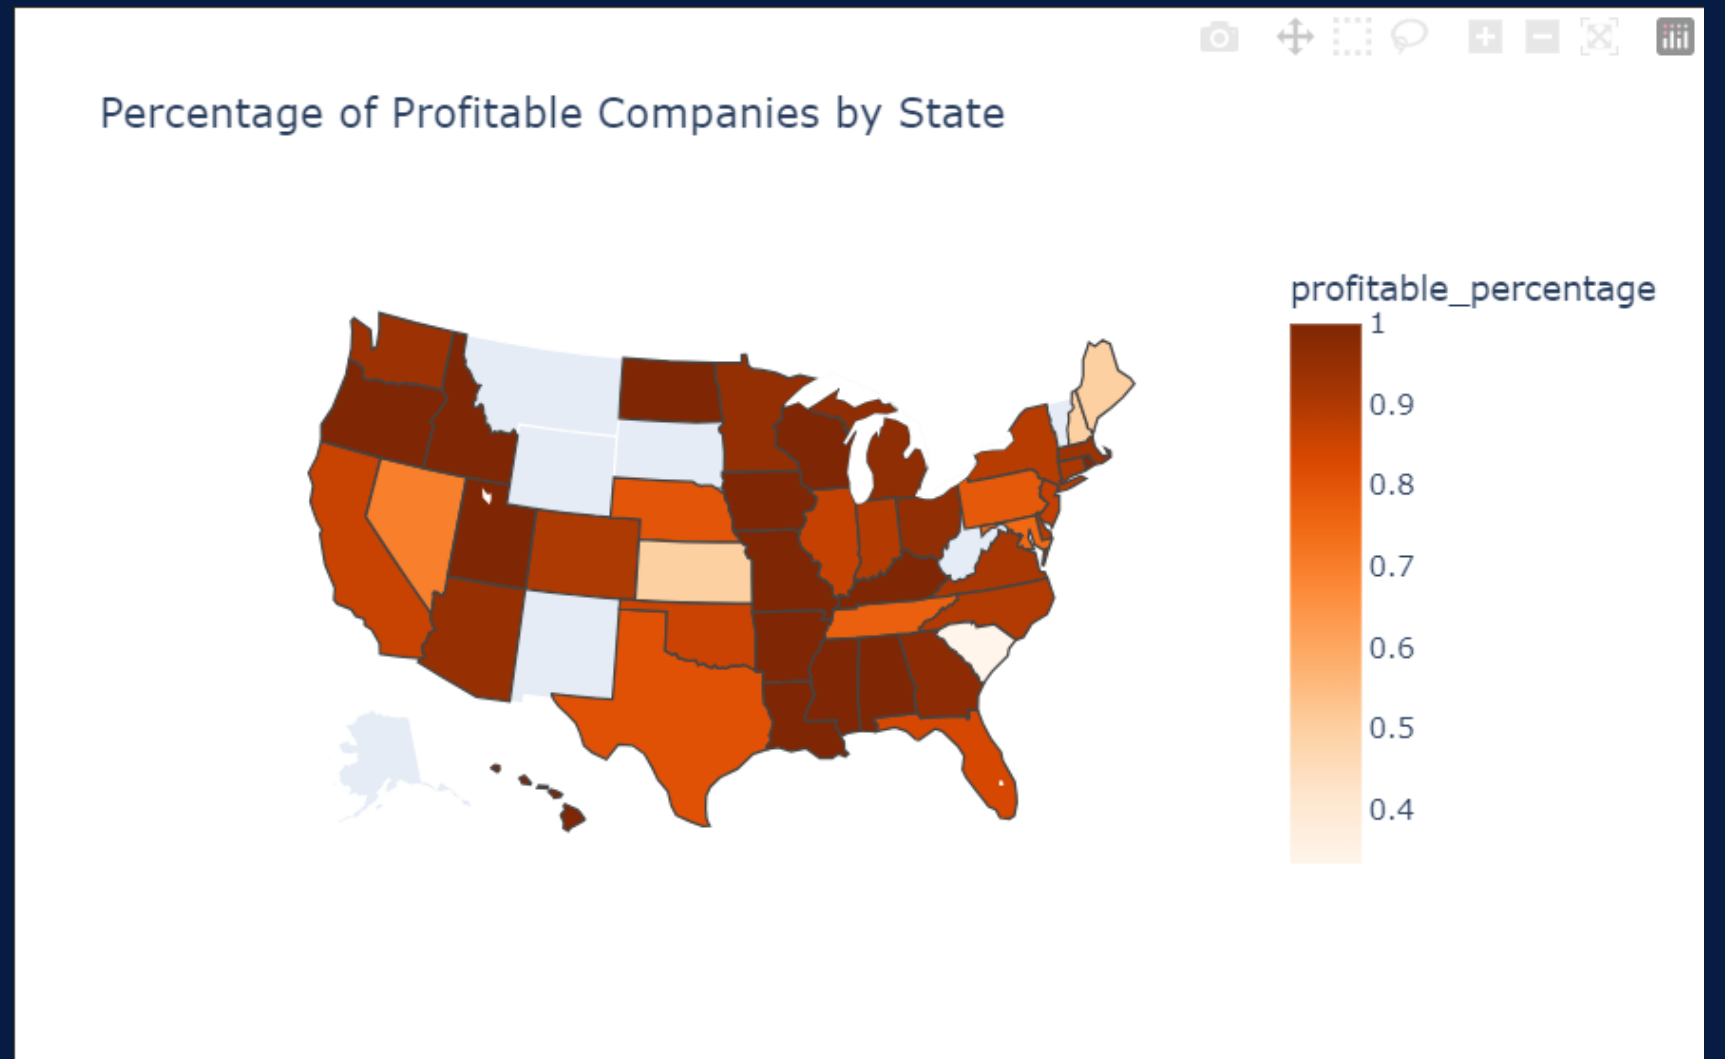
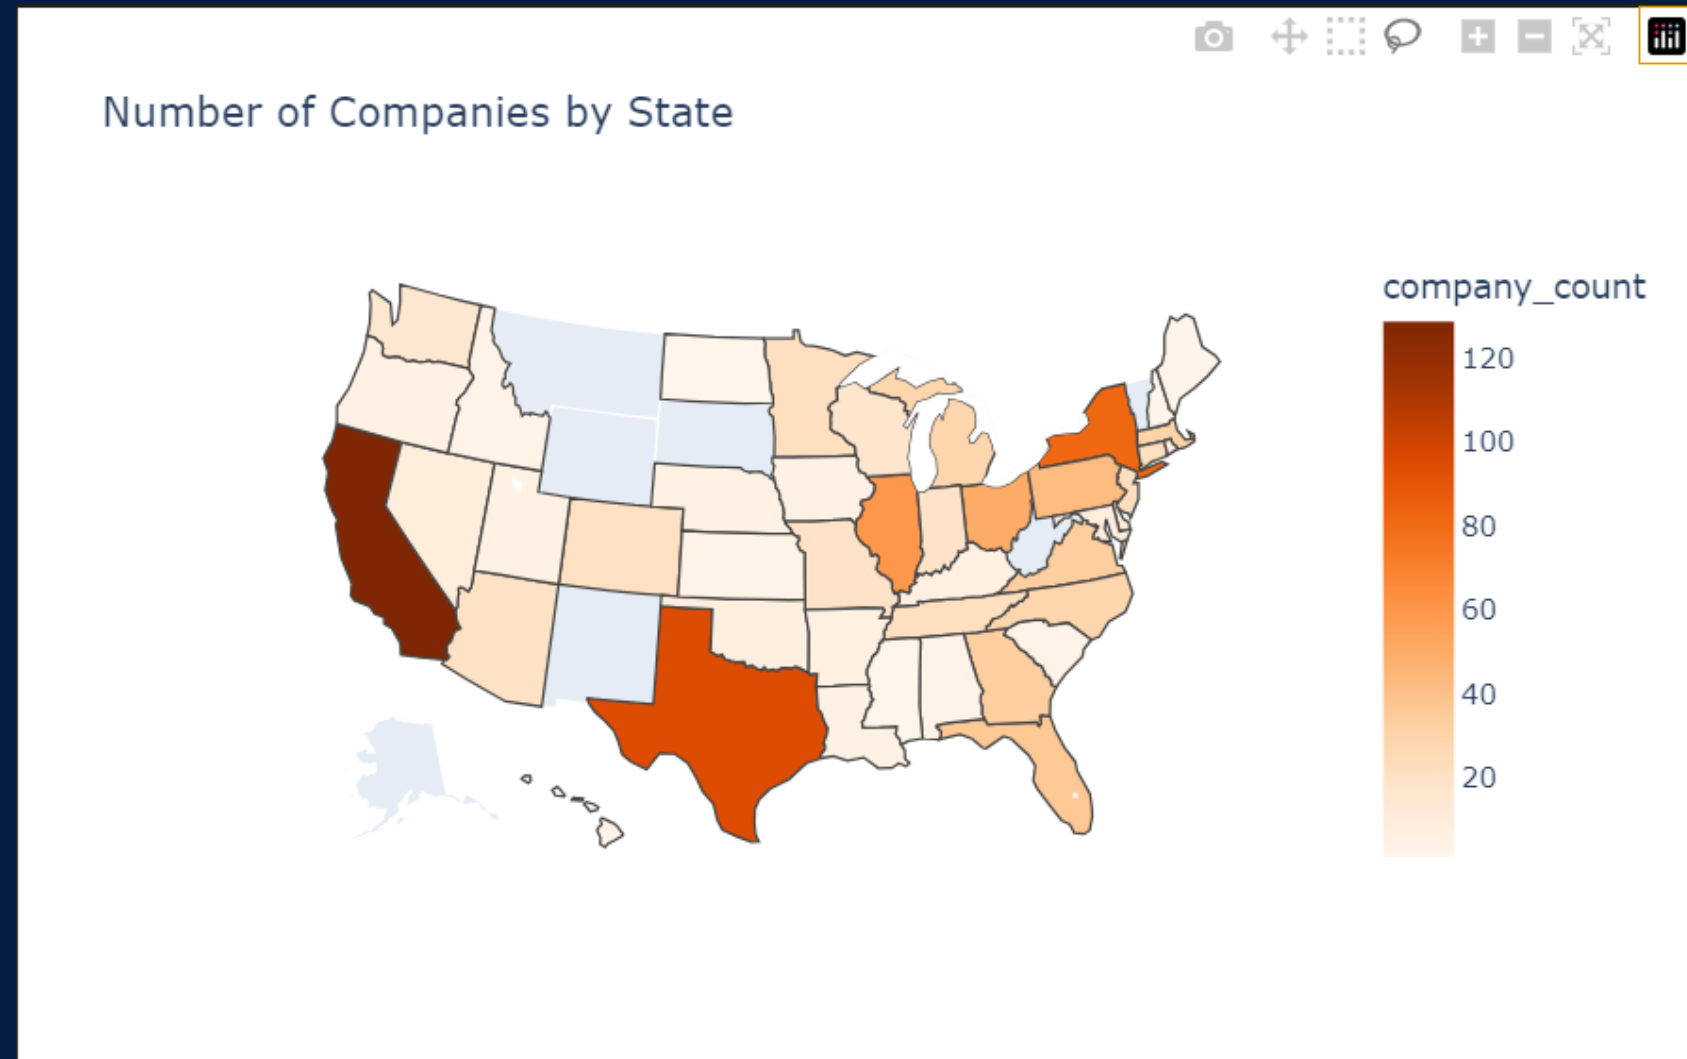
Histograms

As we can see the majority of data is located close to 0, having small "Market Cap", "revenue", "num. of employees" and "profit". This means it does not have a Gaussian distribution and i should normalize it.

Correlation matrix




Looking through cities and states




Exploring sector column

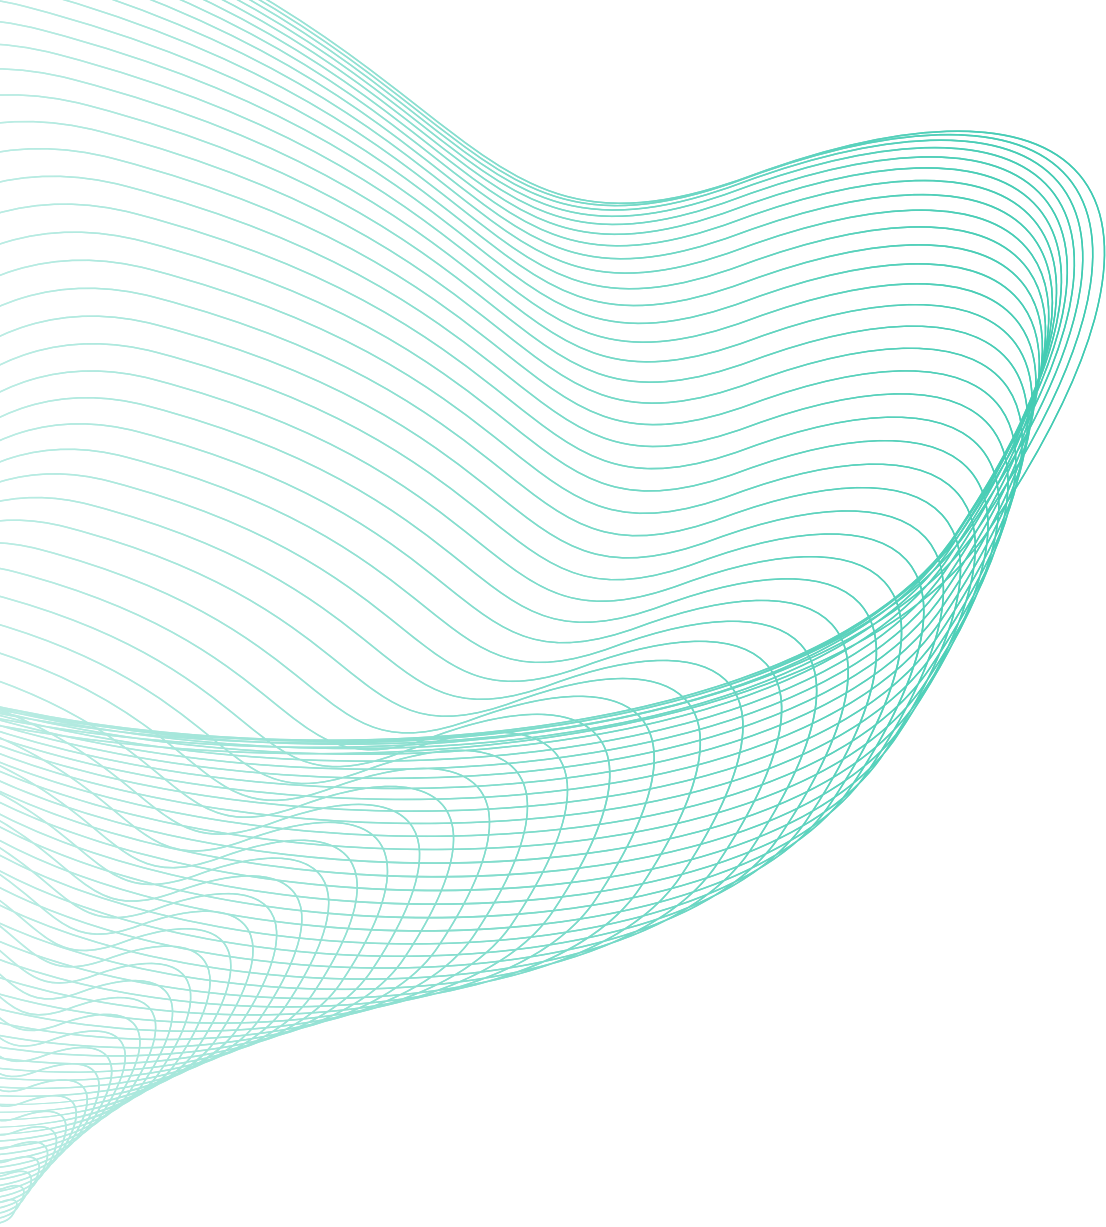
Number of Companies by Sector



Average Market Cap by Sector



Data Preprocessing



1) Dropping unique columns for each company

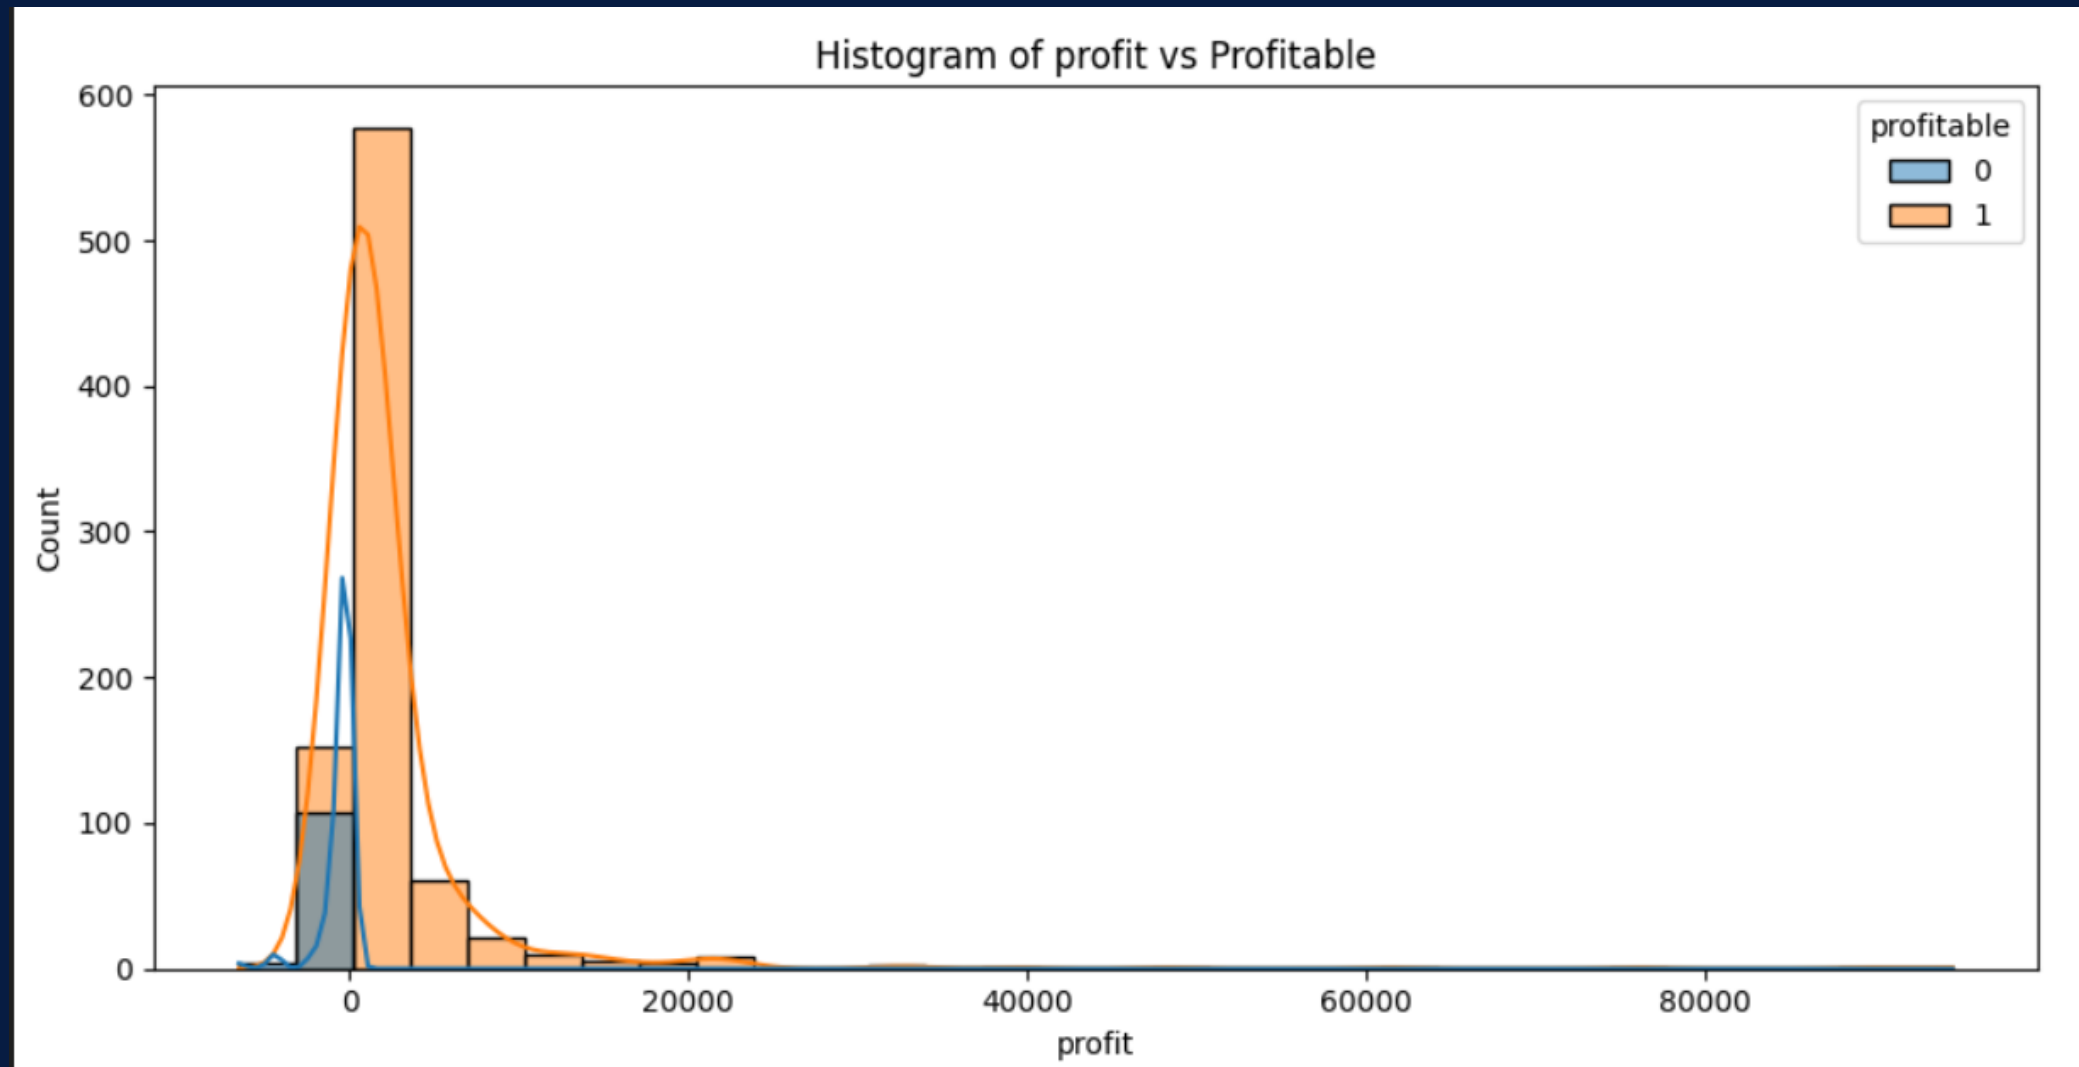
2) Dropping "prev_rank" as half values are missing

```
count_missing = df['prev_rank'].replace(' ', np.nan).astype(float).isna().sum()  
print(f"Missing values in 'prev_rank' column: {count_missing}")
```

```
Missing values in 'prev_rank' column: 531
```

```
df = df.drop(['company', 'CEO', 'Website', 'Ticker',  
             'prev_rank',  
             'profit'], axis=1)
```

3) "profit" column as it directly affects target column



Is profit predicting profitable column ?

if profit <= 0:
 profitable = 0
else:
 profitable = 1

```
# Non-Positive profit and profitable companies
negative_companies = df[df['profit'] <= 0]
negative_profit = negative_companies['profitable'].value_counts()
negative_profit
```

```
profitable
0      111
```

Binary encoding + Missing values

```
columns_to_map = ['newcomer', 'ceo_founder',  
                  'ceo_woman', 'profitable']  
for column in columns_to_map:  
    df[column] = df[column].map({'yes': 1, 'no': 0})
```


rank	0
rank_change	0
revenue	0
num. of employees	1
sector	0
city	0
state	0
newcomer	0
ceo_founder	0
ceo_woman	0
profitable	0
Market Cap	39

```
# Missing value imputation with KNNImputer  
  
imputer = KNNImputer(n_neighbors=5)  
# Getting the columns with missing values  
columns_with_missing_values = df_1.columns[df_1.isna().any()].tolist()  
  
for column in columns_with_missing_values:  
    missing_column = df_1[column]  
    missing_column_2d = missing_column.values.reshape(-1, 1)  
    imputed_column_2d = imputer.fit_transform(missing_column_2d)  
    df_1[column] = imputed_column_2d.flatten()  
  
# Assigning in df_1 the non-numerical columns from df by rank  
df_1 = pd.merge(df_1, concatting_columns_df, on='rank', how='left')  
df_1 = df_1.drop('rank', axis=1)  
df = df_1  
df.info()
```

One Hot 'city' and 'state' columns


Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	revenue	1000 non-null	float64
1	num. of employees	1000 non-null	float64
2	newcomer	1000 non-null	int64
3	ceo_founder	1000 non-null	int64
4	ceo_woman	1000 non-null	int64
5	profitable	1000 non-null	int64
6	Market Cap	1000 non-null	float64
7	sector	1000 non-null	object
8	city	1000 non-null	object
9	state	1000 non-null	object
10	rank_change	1000 non-null	float64



There are 400 cities
There are 46 states
Top states by city count:

state	
CA	50
PA	26
TX	22
IL	22
MI	22
OH	22
NJ	19
FL	19
MA	16
NY	16

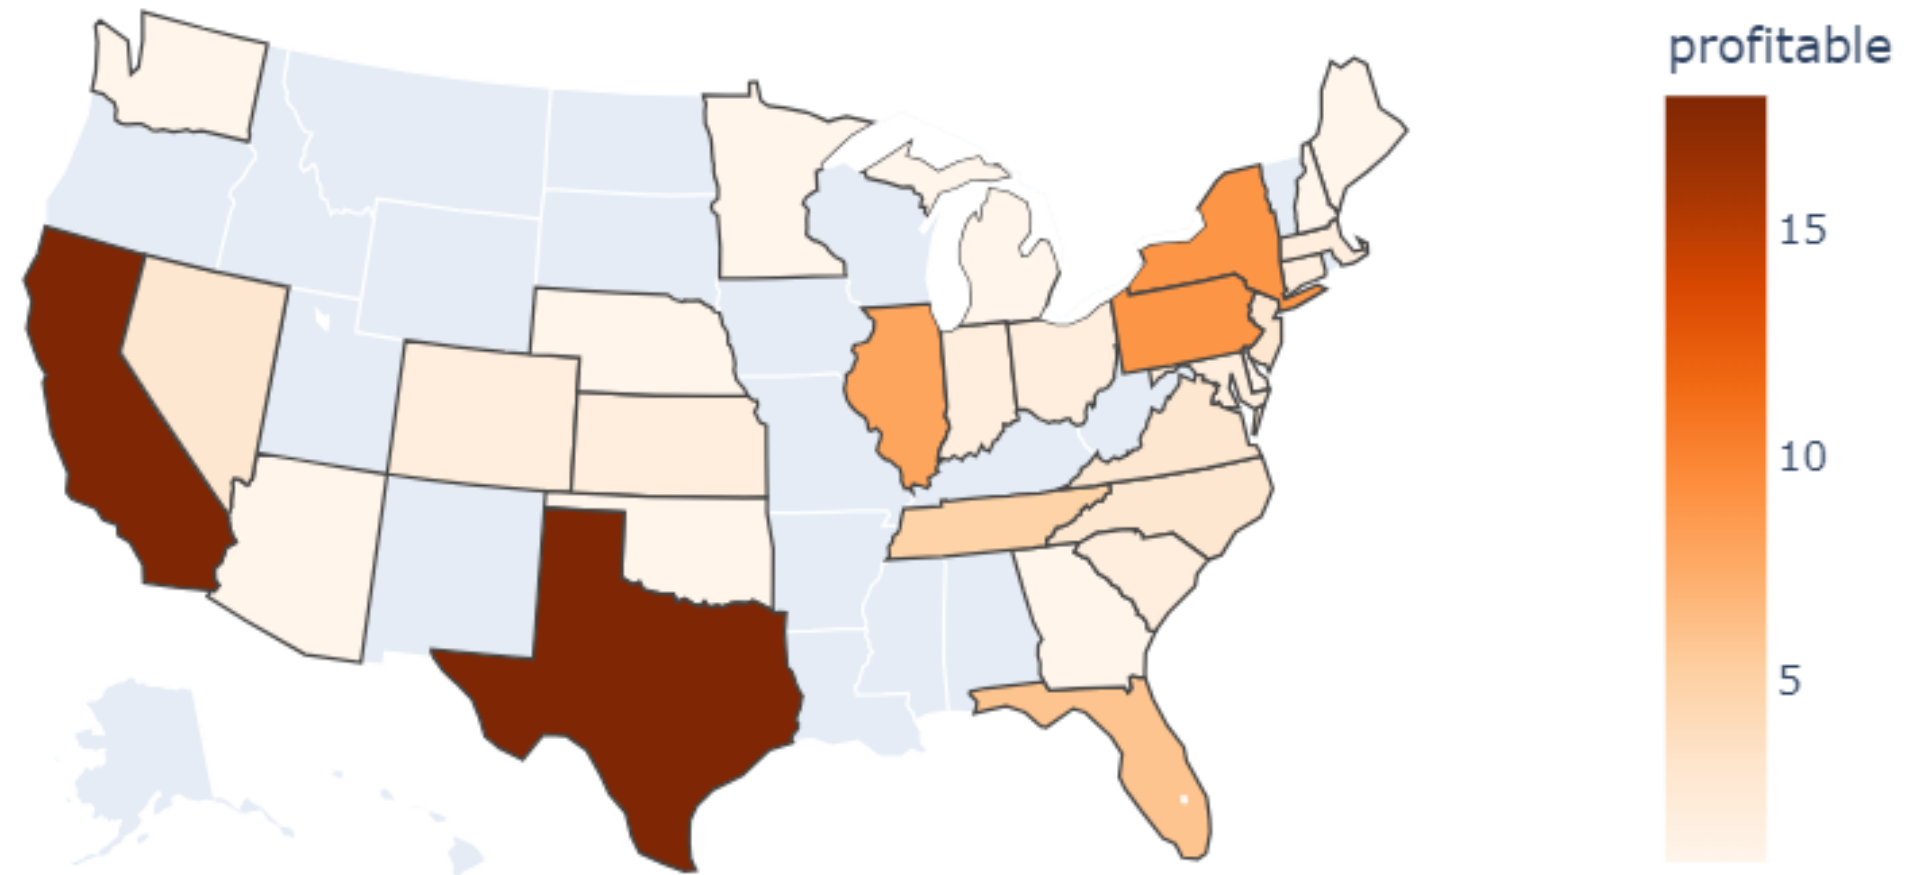


Data columns (total 74 columns):

#	Column	Non-Null Count	Dtype
0	revenue	1000 non-null	float64
1	num. of employees	1000 non-null	float64
2	newcomer	1000 non-null	int64
3	ceo_founder	1000 non-null	int64
4	ceo_woman	1000 non-null	int64
5	profitable	1000 non-null	int64
6	Market Cap	1000 non-null	float64
7	sector	1000 non-null	object
8	rank_change	1000 non-null	float64
9	state_CT	1000 non-null	bool
10	state_FL	1000 non-null	bool
11	state_GA	1000 non-null	bool
12	state_IL	1000 non-null	bool
13	state_MA	1000 non-null	bool
14	state_MI	1000 non-null	bool
15	state_MN	1000 non-null	bool
16	state_NC	1000 non-null	bool
17	state_NJ	1000 non-null	bool
18	state_NY	1000 non-null	bool
19	state_OH	1000 non-null	bool
...			
72	city_Wilmington	1000 non-null	bool
73	city_city_other	1000 non-null	bool

One Hot 'city' and 'state' columns

Unprofitable Companies by State



Same approach for sector column



```
# One hot sector by top 10 and others
top_10 = df['sector'].value_counts().index[:10]
df['sector'] = df['sector'].apply(lambda x: x if x in top_10 else 'others')
df = pd.get_dummies(df, columns=['sector'], drop_first=True)
df.info()
```



Data columns (total 83 columns):

#	Column	Non-Null	Count	Dtype
0	revenue	1000	non-null	float64
1	num. of employees	1000	non-null	float64
2	newcomer	1000	non-null	int64
3	ceo_founder	1000	non-null	int64
4	ceo_woman	1000	non-null	int64
5	profitable	1000	non-null	int64
6	Market Cap	1000	non-null	float64
7	rank_change	1000	non-null	float64
8	state_CT	1000	non-null	bool
9	state_FL	1000	non-null	bool
10	state_GA	1000	non-null	bool
11	state_IL	1000	non-null	bool
12	state_MA	1000	non-null	bool
13	state_MI	1000	non-null	bool
14	state_MN	1000	non-null	bool
15	state_NC	1000	non-null	bool
16	state_NJ	1000	non-null	bool
17	state_NY	1000	non-null	bool
18	state_OH	1000	non-null	bool
19	state_PA	1000	non-null	bool
...				
81	sector_Transportation	1000	non-null	bool
82	sector_others	1000	non-null	bool

Choosing best number of features



Data columns (total 67 columns):

#	Column	Non-Null Count	Dtype
0	revenue	1000 non-null	float64
1	num. of employees	1000 non-null	float64
2	newcomer	1000 non-null	int64
3	ceo_founder	1000 non-null	int64
4	ceo_woman	1000 non-null	int64
5	Market Cap	1000 non-null	float64
6	rank_change	1000 non-null	float64
7	state_FL	1000 non-null	bool
8	state_GA	1000 non-null	bool
9	state_MA	1000 non-null	bool
10	state_MI	1000 non-null	bool
11	state_MN	1000 non-null	bool
12	state_NY	1000 non-null	bool
13	state_OH	1000 non-null	bool
14	state_PA	1000 non-null	bool
15	state_TX	1000 non-null	bool
16	state_VA	1000 non-null	bool
17	state_state_other	1000 non-null	bool
18	city_Atlanta	1000 non-null	bool
19	city_Austin	1000 non-null	bool
...			
65	sector_others	1000 non-null	bool
66	profitable	1000 non-null	int64


Removing 3% of outliers from train set

```
# Removing outliers with Isolation Forest
outlier_detector = IsolationForest(contamination=0.03)
outlier_detector.fit(X_train)
no_outliers = outlier_detector.predict(X_train)
no_outliers = no_outliers == 1
X_train, y_train = X_train[no_outliers], y_train[no_outliers]
print("Training set - X:", X_train.shape, "y:", y_train.shape)
```


```
Training set - X: (776, 66) y: (776,)
```

SMOTEENN

imblearn.combine




CRUCIO




Data Processing

Random Forest



```
Accuracy: 0.88  
Precision: 0.8974358974358975  
Recall: 0.9776536312849162  
F1: 0.9358288770053476  
[[ 1 20]  
 [ 4 175]]
```



Feature Importances




SHAP Plots




XGBoost

```
Accuracy: 0.77  
Precision: 0.9182389937106918  
Recall: 0.8156424581005587  
F1: 0.8639053254437871  
[[ 8 13]  
 [33 146]]
```




DecisionTreeClassifier

```
Accuracy: 0.84  
Precision: 0.9152542372881356  
Recall: 0.9050279329608939  
F1: 0.9101123595505618  
[[ 6 15]  
 [17 162]]
```



Tree visualization




Scaling data

```
#Normalize the data
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```


Logistic Regression

```
Accuracy: 0.82  
Precision: 0.9181286549707602  
Recall: 0.8770949720670391  
F1: 0.8971428571428571  
[[ 7 14]  
 [22 157]]
```




AdaBoost

```
Accuracy: 0.855  
Precision: 0.9166666666666666  
Recall: 0.9217877094972067  
F1: 0.9192200557103063  
[[ 6 15]  
 [14 165]]
```




KNN

```
Accuracy: 0.75  
Precision: 0.910828025477707  
Recall: 0.7988826815642458  
F1: 0.8511904761904762  
[[ 7 14]  
 [36 143]]
```



Models performance



**THANK YOU
FOR YOUR
ATTENTION**

